

---

---

# Distributed Computing

## Module 1 -Lecture 5/6

Maresh C.

Centre for High Performance Computing  
FISAT

---

# Logical vs. Physical Concurrency

- Two events are logically concurrent if and only if they do not causally affect each other.
- Physical concurrency has a connotation that the events occur at the same instant in physical time.
- Two or more events may be logically concurrent even though they do not occur at the same instant in physical time.

# Models of Communication Network

- FIFO Model
- non-FIFO model
- Causal ordering Model ( CO)

CO: For any two messages  $m_{ij}$  and  $m_{kj}$ , if  $\text{send}(m_{ij}) \rightarrow \text{send}(m_{kj})$ ,  
then  $\text{rec}(m_{ij}) \rightarrow \text{rec}(m_{kj})$

# Global State of a Distributed System

- The global state of a distributed system is a collection of the local states of the processes and the channels.

## Process State

Let  $LS_i^x$  denote the state of process  $p_i$  after the occurrence of event  $e_i^x$  and before the event  $e_i^{x+1}$ .  $LS_i^0$  denotes the initial state of process  $p_i$ .  $LS_i^x$  is a result of the execution of all the events executed by process  $p_i$  till  $e_i^x$ . Let  $send(m) \leq LS_i^x$  denote the fact that  $\exists y: 1 \leq y \leq x :: e_i^y = send(m)$ . Likewise, let  $rec(m) \not\leq LS_i^x$  denote the fact that  $\forall y: 1 \leq y \leq x :: e_i^y \neq rec(m)$ .

# Global State of a Distributed System

The state of a channel is difficult to state formally because a channel is a distributed entity and its state depends upon the states of the processes it connects. Let  $SC_{ij}^{x,y}$  denote the state of a channel  $C_{ij}$  defined as follows:

$$SC_{ij}^{x,y} = \{m_{ij} \mid send(m_{ij}) \leq e_i^x \wedge rec(m_{ij}) \not\leq e_j^y\}$$

Thus, channel state  $SC_{ij}^{x,y}$  denotes all messages that  $p_i$  sent upto event  $e_i^x$  and which process  $p_j$  had not received until event  $e_j^y$ .


# Global State of a Distributed System

The global state of a distributed system is a collection of the local states of the processes and the channels. Notationally, global state GS is defined as

$$GS = \{ \bigcup_i LS_i^{x_i}, \bigcup_{j,k} SC_{jk}^{y_j, z_k} \}$$


# Global State - consistent

a state will be meaningful provided every message that is recorded as received is also recorded as sent.



# Global State - consistent or not


$\{LS_1^1, \overline{LS_2^3}, \overline{LS_3^3}, LS_4^2\}$





# Global State - consistent or not

$\{LS_1^2, LS_2^4, LS_3^4, LS_4^2\}$




# Cuts of a Distributed Computation

**“In the space-time diagram of a distributed computation, a cut is a zigzag line joining one arbitrary point on each process line.”**

- **A cut slices the space-time diagram, and thus the set of events in the distributed computation, into a PAST and a FUTURE.**
- **The PAST contains all the events to the left of the cut and the FUTURE contains all the events to the right of the cut**
- **For a cut  $C$ , let  $PAST(C)$  and  $FUTURE(C)$  denote the set of events in the PAST and FUTURE of  $C$ , respectively**
- **Every cut corresponds to a global state and every global state can be graphically represented as a cut in the computation space-time diagram.**
- **Cuts in a space-time diagram provide a powerful graphical aid in representing and reasoning about global states of a computation.**

# Cuts of a Distributed Computation



# Cuts of a Distributed Computation

**“In the space-time diagram of a distributed computation, a cut is a zigzag line joining one arbitrary point on each process line.”**

- **In a consistent cut, every message received in the PAST of the cut was sent in the PAST of that cut. (In, cut C2 is a consistent cut.)**
- **All messages that cross the cut from the PAST to the FUTURE are in transit in the corresponding consistent global state.**
- **A cut is inconsistent if a message crosses the cut from the FUTURE to the PAST. (In Figure, cut C1 is an inconsistent cut.)**


# Past and Future Cones of an Event

## “Past Cone of an Event”

- An event  $e_j$  could have been affected only by all events  $e_i$  such that  $e_i \rightarrow e_j$ .
- In this situation, all the information available at  $e_i$  could be made accessible at  $e_j$ .
- All such events  $e_i$  belong to the past of  $e_j$ . Let  $Past(e_j)$  denote all events in the past of  $e_j$  in a computation  $(H, \rightarrow)$ . Then

$$Past(e_j) = \{e_i | \forall e_i \in H, e_i \rightarrow e_j\}.$$

# Past and Future Cones of an Event



# Past and Future Cones of an Event

## “Future Cone of an Event”

- The future of an event  $e_j$ , denoted by  $\text{Future}(e_j)$ , contains all events  $e_i$  that are causally affected by  $e_j$
- In a computation  $\langle H, \rightarrow \rangle$ ,  $\text{Future}(e_j)$  is defined as:

$$\text{Future}(e_j) = \{e_i \mid \forall e_i \in H, e_j \rightarrow e_i\}.$$

# Models of Process Communications

- There are two basic models of process communications – synchronous and asynchronous.
- The synchronous communication model is a blocking type where on a message send, the sender process blocks until the message has been received by the receiver process.
- The sender process resumes execution only after it learns that the receiver process has accepted the message.
- Thus, the sender and the receiver processes must synchronize to exchange a message.



# Models of Process Communications

- **Asynchronous communication model is a non-blocking type where the sender and the receiver do not synchronize to exchange a message.**
- **After having sent a message, the sender process does not wait for the message to be delivered to the receiver process.**
- **The message is buffered by the system and is delivered to the receiver process when it is ready to accept the message.**

# Models of Process Communications

- **Asynchronous communication provides higher parallelism because the sender process can execute while the message is in transit to the receiver.**
- **However, A buffer overflow may occur if a process sends a large number of messages in a burst to another process.**
- **An implementation of asynchronous communication requires more complex buffer management. In addition, due to higher degree of parallelism and non-determinism, it is much more difficult to design, verify, and implement distributed algorithms for asynchronous communications.**
- **Synchronous communication is simpler to handle and implement. However, due to frequent blocking, it is likely to have poor performance and is likely to be more prone to deadlocks.**